Which parser for which document

Every kind of document situation we tested, from plain text to a table inside a table and noisy scans: the parsers that did best and the ones to avoid. 25 situations, 62 + 42 test documents. In 8 situations the app's first pick is not the best.

1. Cheat sheet

Situation	Best parsers	Weak or fails	App picks (by file type)	Is that pick good here?
Text documents				
Plain text pages	PyMuPDF, hybrid_pdf, smart_pdf	Unstructured, MarkItDown, Tesseract OCR, pdfminer	PyMuPDF	yes
Long documents with nested numbering	ADI + vision (paid)	smart_pdf, PyMuPDF, hybrid_pdf, pdfplumber	OpenDataLoader	no
Tables				
Simple clean tables	Docling, PyMuPDF4LLM, OpenDataLoader, PyMuPDF, hybrid_pdf, smart_pdf	MarkItDown, Unstructured, LangChain / LlamaIndex / Haystack loaders, pdfminer	OpenDataLoader	yes
One table over many pages	OpenDataLoader, PyMuPDF4LLM	MarkItDown, Docling, smart_pdf, Unstructured	OpenDataLoader	yes
Very wide table (12 columns)	Docling, OpenDataLoader, PyMuPDF, hybrid_pdf, smart_pdf	MarkItDown, Unstructured, Tesseract OCR	OpenDataLoader	yes
Merged cells and two-row headers	PyMuPDF4LLM, Docling, ADI + vision (paid)	MarkItDown, smart_pdf, Unstructured, Tesseract OCR	OpenDataLoader	no
A table inside a table	OpenDataLoader, PyMuPDF4LLM	ADI + vision (paid), Tesseract OCR, Docling, Unstructured	OpenDataLoader	yes
Two tables side by side	Docling, PyMuPDF4LLM	PyMuPDF, hybrid_pdf, smart_pdf, OpenDataLoader	OpenDataLoader	no
Forms with boxes and check marks	PyMuPDF, hybrid_pdf, smart_pdf, OpenDataLoader	Docling	OpenDataLoader	yes
Invoices and orders with logo, QR code, stamp	ADI + vision (paid)	pdfplumber, PyMuPDF, hybrid_pdf, MarkItDown	OpenDataLoader	no
Charts and pictures				
Charts with numbers	smart_pdf	PyMuPDF, hybrid_pdf, pdfplumber, MarkItDown	smart_pdf	yes
Photos, diagrams, brochures, slides (no chart values)	OpenDataLoader, PyMuPDF4LLM	Unstructured, MarkItDown, LangChain / LlamaIndex / Haystack loaders, Tesseract OCR	smart_pdf	no
Two-column layout	smart_pdf	pdfplumber, LangChain / LlamaIndex / Haystack loaders, EasyOCR, Tesseract OCR	OpenDataLoader	no
Dense financial research reports (9 to 66 pages)	Azure Document Intelligence (paid), ADI + vision (paid), PyMuPDF4LLM, OpenDataLoader	EasyOCR	smart_pdf	no
Long documents that mix everything	Docling	hybrid_pdf, pdfplumber, PyMuPDF, MarkItDown	smart_pdf	no
Scans and images				
Clean scans	Docling	EasyOCR, hybrid_pdf, Tesseract OCR, smart_pdf	Docling	yes
Noisy, tilted or low-quality scans, phone photos	Docling	Unstructured, EasyOCR, hybrid_pdf, Tesseract OCR	Docling	yes

Long scan (31 pages)	Docling	PyMuPDF4LLM, Unstructured, hybrid_pdf, Tesseract OCR	Docling	yes
Pages with text and pages that are scans	Docling, PyMuPDF4LLM	EasyOCR, Tesseract OCR, PyMuPDF, OpenDataLoader	PyMuPDF4LLM	yes
Image files (png, jpg)	Docling	-	Docling	yes
Office files				
Word	Unstructured	-	native Word reader	native, on purpose
PowerPoint	Docling	Unstructured	native PowerPoint reader	native, on purpose
Excel	native Excel reader, MarkItDown	-	native Excel reader	yes
CSV	native Excel reader, MarkItDown, Unstructured	-	native Excel reader	yes
Plain text and Markdown	plain text reader	Unstructured	plain text reader	yes

Best (equal) = within 0.1 of the top score; those parsers are tied. **Also good** = within 0.6. **Weak** = lower. **Fails** = average below 1 (returns nothing or crashes on most files). Scores are the average over the test documents of that situation, 0 to 5.

2. Details for each situation

Text documents

Plain text pages

Reports, contracts, letters, resumes, meeting minutes.

Level	Parsers (average score, 0 to 5)
Best (equal)	PyMuPDF 4.92, hybrid_pdf 4.92, smart_pdf 4.92
Also good	Docling 4.82, OpenDataLoader 4.70, LangChain / LlamaIndex / Haystack loaders 4.65, PyMuPDF4LLM 4.60, pdfplumber 4.48
Weak	Unstructured 4.17, MarkItDown 4.06, Tesseract OCR 3.81, pdfminer 3.46, EasyOCR 2.63

Tip: Every parser with a text layer does well. Pick the fastest.

App pick: OK here. The app picks PyMuPDF for this type of file: PyMuPDF scores 4.92, within 0.1 of the best.

Tested on 6 documents: AI_agent_p1-3, contract, resume, multilingual, rotated_pages, board_minutes.

Long documents with nested numbering

Contracts and reports of 4 to 46 pages, clauses like 1.2.1, header and footer on every page.

Level	Parsers (average score, 0 to 5)
Best (equal)	ADI + vision (paid) 4.90
Also good	Docling 4.54, PyMuPDF4LLM 4.48, OpenDataLoader 4.43
Weak	smart_pdf 4.27, PyMuPDF 4.26, hybrid_pdf 4.26, pdfplumber 4.11, Unstructured 3.88, MarkItDown 3.83, Tesseract OCR 3.80

Tip: Length is not a problem for the fast parsers. Docling and PyMuPDF4LLM lead; Docling takes minutes on the long ones.

App pick: NOT the best here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.43, 0.47 below the best (4.90).

Tested on 4 documents: contract_nested_clauses, long_annual_report_20p, contract_45p, 03_legal_contract.

Tables

Simple clean tables

One normal grid: balance sheet, invoice, datasheet, lab report, purchase order.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.84, PyMuPDF4LLM 4.81, OpenDataLoader 4.80, PyMuPDF 4.75, hybrid_pdf 4.75, smart_pdf 4.75
Also good	pdfplumber 4.41
Weak	MarkItDown 4.21, Unstructured 3.75, LangChain / LlamaIndex / Haystack loaders 3.74, pdfminer 3.65, Tesseract OCR 3.29, EasyOCR 1.86

Tip: Most layout-aware parsers are good. Plain text-only readers lose the structure.

App pick: OK here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.80, within 0.1 of the best.

Tested on 8 documents: balance_sheet, invoice, digital, datasheet, bank_statement, lab_report +2 more.

One table over many pages

A ledger or statement that continues for 3 to 100 pages, header repeated on each page.

Level	Parsers (average score, 0 to 5)
Best (equal)	OpenDataLoader 4.93, PyMuPDF4LLM 4.91
Also good	pdfplumber 4.54, PyMuPDF 4.38, hybrid_pdf 4.38
Weak	MarkItDown 3.77, Docling 3.67, smart_pdf 3.63, Unstructured 3.60, Tesseract OCR 2.77

Tip: Pages do not break the good parsers. PyMuPDF and hybrid_pdf are much lower on the 100-page ledger; Docling and smart_pdf did not finish it.

App pick: OK here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.93, within 0.1 of the best.

Tested on 4 documents: multipage_table, ledger_repeated_header, bank_statement, transaction_ledger_100p.

Very wide table (12 columns)

Bill of materials in small print.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 5.00, OpenDataLoader 4.99, PyMuPDF 4.93, hybrid_pdf 4.93, smart_pdf 4.93
Also good	PyMuPDF4LLM 4.77, ADI + vision (paid) 4.58, pdfplumber 4.53
Weak	MarkItDown 3.54, Unstructured 3.54, Tesseract OCR 2.90

Tip: The top parsers keep every column. MarkItDown and Unstructured lose columns.

App pick: OK here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.99, within 0.1 of the best.

Tested on 1 document (one document only: a hint, not proof): bom_wide_table.

Merged cells and two-row headers

Financial statements, clinical results, bill of lading, requirement matrix with spanning cells.

Level	Parsers (average score, 0 to 5)
Best (equal)	PyMuPDF4LLM 4.77, Docling 4.72, ADI + vision (paid) 4.69
Also good	OpenDataLoader 4.60, pdfplumber 4.25, PyMuPDF 4.23, hybrid_pdf 4.23
Weak	MarkItDown 3.92, smart_pdf 3.82, Unstructured 3.57, Tesseract OCR 3.16

Tip: Use PyMuPDF4LLM, OpenDataLoader or Docling. PyMuPDF, hybrid_pdf and smart_pdf break the table apart or repeat the text.

App pick: NOT the best here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.60, 0.16 below the best (4.77).

Tested on 5 documents: financial_statement_multiheader, clinical_trial_table, annual_report_income_statement, bill_of_lading, rfp_requirements_matrix.

A table inside a table

Policy schedule: the coverage cell holds its own table.

Level	Parsers (average score, 0 to 5)
Best (equal)	OpenDataLoader 4.77, PyMuPDF4LLM 4.77
Also good	MarkItDown 4.58, pdfplumber 4.36
Weak	ADI + vision (paid) 3.89, Tesseract OCR 3.67, Docling 3.62, Unstructured 3.58, PyMuPDF 2.92, hybrid_pdf 2.92, smart_pdf 2.92

Tip: Only PyMuPDF4LLM and OpenDataLoader kept it right. Docling is weaker. One document only, so treat this as a hint.

App pick: OK here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.77, within 0.1 of the best.

Tested on 1 document (one document only: a hint, not proof): policy_schedule_nested.

Two tables side by side

Payslip: earnings on the left, deductions on the right.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.69, PyMuPDF4LLM 4.62
Also good	MarkItDown 4.34
Weak	PyMuPDF 4.05, hybrid_pdf 4.05, smart_pdf 4.05, OpenDataLoader 3.94, pdfplumber 3.67, Unstructured 3.61, Tesseract OCR 3.53

Tip: Docling and PyMuPDF4LLM are best. OpenDataLoader, the app's usual table parser, is clearly lower here. One document only.

App pick: NOT the best here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 3.94, 0.75 below the best (4.69).

Tested on 1 document (one document only: a hint, not proof): payslip.

Forms with boxes and check marks

Compliance form, insurance claim, discharge summary, fillable form.

Level	Parsers (average score, 0 to 5)
Best (equal)	PyMuPDF 3.99, hybrid_pdf 3.99, smart_pdf 3.99, OpenDataLoader 3.98
Also good	PyMuPDF4LLM 3.85, Unstructured 3.68, pdfplumber 3.67, EasyOCR 3.66, Tesseract OCR 3.65, LangChain / LlamaIndex / Haystack loaders 3.64, MarkItDown 3.63, pdfminer 3.61
Weak	Docling 2.77

Tip: The PyMuPDF family and OpenDataLoader are fine. Check marks and boxes are easy.

App pick: OK here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 3.98, within 0.1 of the best.

Tested on 4 documents: form_grid_ac4, 06_insurance_claim_form, 05_healthcare_discharge_summary, form_fields.

Invoices and orders with logo, QR code, stamp

Tax invoice, purchase order, manifest.

Level	Parsers (average score, 0 to 5)
Best (equal)	ADI + vision (paid) 4.68
Also good	OpenDataLoader 4.46, Docling 4.42, PyMuPDF4LLM 4.34
Weak	pdfplumber 3.97, PyMuPDF 3.92, hybrid_pdf 3.92, MarkItDown 3.88, smart_pdf 3.88, EasyOCR 3.68, Unstructured 3.62, LangChain / LlamaIndex / Haystack loaders 3.57, pdfminer 3.35, Tesseract OCR 3.05

Tip: OpenDataLoader is best on the complex invoice. smart_pdf adds a description of the QR code that is not in the answer.

App pick: NOT the best here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.46, 0.22 below the best (4.68).

Tested on 3 documents: invoice_complex, purchase_order, 07_logistics_shipping_manifest.

Charts and pictures

Charts with numbers

Annual reports, utility bill, research paper with a bar chart or pie chart.

Level	Parsers (average score, 0 to 5)
Best (equal)	smart_pdf 3.82
Also good	PyMuPDF4LLM 3.38, OpenDataLoader 3.24
Weak	PyMuPDF 3.09, hybrid_pdf 3.09, pdfplumber 3.07, MarkItDown 2.79, Docling 2.78, Unstructured 2.63, LangChain / LlamaIndex / Haystack loaders 2.57, EasyOCR 2.55, Tesseract OCR 2.54, pdfminer 2.50

Tip: Only smart_pdf reads the numbers inside a chart (it uses a vision model). Every other parser reads only the words printed on the chart.

App pick: OK here. The app picks smart_pdf for this type of file: smart_pdf scores 3.82, within 0.1 of the best.

Tested on 10 documents: chart_report, chart_sample, annual_report_page, 01_corporate_annual_report, 04_scientific_research_paper, annual_report_spread +4 more.

Photos, diagrams, brochures, slides (no chart values)

Datasheet with a photo, flowchart, brochure, slides saved as PDF, web pages.

Level	Parsers (average score, 0 to 5)
Best (equal)	OpenDataLoader 4.19, PyMuPDF4LLM 4.14
Also good	Docling 3.99, PyMuPDF 3.98, hybrid_pdf 3.98, pdfplumber 3.87, smart_pdf 3.60
Weak	Unstructured 3.53, MarkItDown 3.49, LangChain / LlamaIndex / Haystack loaders 3.48, Tesseract OCR 3.08, pdfminer 2.96, EasyOCR 2.26

Tip: Free parsers (PyMuPDF4LLM) can score higher than smart_pdf here: the vision model adds descriptions that are not in the answer.

App pick: NOT the best here. The app picks smart_pdf for this type of file: smart_pdf scores 3.60, 0.60 below the best (4.19).

Tested on 8 documents: datasheet_with_photo, flowchart_doc, product_brochure, slides, 10_mixed_parser_killer, PIB_home_p1-6 +2 more.

Two-column layout

Research paper with two columns and a table.

Level	Parsers (average score, 0 to 5)
Best (equal)	smart_pdf 4.52
Also good	PyMuPDF4LLM 4.37, Docling 4.37, OpenDataLoader 4.34, PyMuPDF 4.05, hybrid_pdf 4.05
Weak	pdfplumber 3.79, LangChain / LlamaIndex / Haystack loaders 3.71, EasyOCR 3.62, Tesseract OCR 3.09, Unstructured 2.72, MarkItDown 2.67
Fails / empty	pdfminer 0.00

Tip: Reading order matters. smart_pdf and the layout-aware parsers do best.

App pick: NOT the best here. The app picks OpenDataLoader for this type of file: OpenDataLoader scores 4.34, 0.18 below the best (4.52).

Tested on 2 documents: two_column_paper, research_two_column.

Dense financial research reports (9 to 66 pages)

Broker and market-research PDFs with many wide tables (some on rotated pages), charts and small text boxes.

Level	Parsers (average score, 0 to 5)
Best (equal)	Azure Document Intelligence (paid) 1.64, ADI + vision (paid) 1.62, PyMuPDF4LLM 1.58, OpenDataLoader 1.56

Also good	pdfplumber 1.51, MarkItDown 1.45, smart_pdf 1.40, PyMuPDF 1.38, hybrid_pdf 1.38, pdfminer 1.31, Tesseract OCR 1.29, haystack_converter 1.29, langchain_loader 1.29, llamaindex_reader 1.29, Unstructured 1.11, Docling 1.09
Fails / empty	EasyOCR 0.70

Tip: Most parsers are equal (within 0.1); OpenDataLoader, pdfplumber, PyMuPDF4LLM and MarkItDown are the fast ones. Docling takes minutes and did not finish the 61 and 66 page files in 25 minutes; EasyOCR did not finish 3 of 6. Azure DI (paid) is first on one report only. smart_pdf does not help here. Answer files cover key pages only, so use the scores for comparing parsers, not as absolute values.

App pick: NOT the best here. The app picks smart_pdf for this type of file: smart_pdf scores 1.40, 0.24 below the best (1.64).

Tested on 6 documents: bernstein_novo_66p, tbrc_mash_market_61p, piper_madrigal_10p, nh_bio_sector_14p, hcw_madrigal_10p, alphasense_mash_report_9p.

Long documents that mix everything

31 to 99 pages with text, merged-cell tables, side-by-side tables, nested tables, charts, photos, two columns and scanned pages in one file.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.37
Also good	PyMuPDF4LLM 4.16, OpenDataLoader 3.83
Weak	hybrid_pdf 3.70, pdfplumber 3.68, PyMuPDF 3.65, MarkItDown 3.56, Unstructured 2.98, Tesseract OCR 2.88, smart_pdf 2.44

Tip: On the two shorter files Docling and PyMuPDF4LLM lead (4.07 to 4.66). Tables are the weak spot: OpenDataLoader and PyMuPDF score only 2.6 to 2.9 on tables. smart_pdf did not finish the 36-page file with scanned pages in 10 minutes and scored lower on the other (photo descriptions). By part did not finish in 15 minutes on the file with scans. On the 99-page file (tested without Docling and Tesseract OCR, too slow) PyMuPDF4LLM was best (3.93) and smart_pdf read only 8 of about 27 charts (cap of 8 per document).

App pick: NOT the best here. The app picks smart_pdf for this type of file: smart_pdf scores 2.44, 1.93 below the best (4.37).

Tested on 3 documents: operations_report_36p, tender_pack_30p, group_report_100p.

Scans and images

Clean scans

Scanned invoice, report page, maintenance order.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.47
Also good	PyMuPDF4LLM 4.36, Unstructured 3.91
Weak	EasyOCR 3.87, hybrid_pdf 3.69, Tesseract OCR 3.69, smart_pdf 3.69
Fails / empty	PyMuPDF 0.00, pdfminer 0.00, pdfplumber 0.00, MarkItDown 0.00, LangChain / LlamaIndex / Haystack loaders 0.00, OpenDataLoader 0.00

Tip: Docling is best. PyMuPDF and the text-only parsers return nothing; Tesseract alone loses tables.

App pick: OK here. The app picks Docling for this type of file: Docling scores 4.47, within 0.1 of the best.

Tested on 4 documents: AI_agent_scanned_clean, scanned_invoice_clean, scanned_report_page, 09_scanned_maintenance_order.

Noisy, tilted or low-quality scans, phone photos

Skewed, blurred, grainy scans and a phone photo of a purchase order.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 3.89
Also good	PyMuPDF4LLM 3.48

Weak	Unstructured 2.73, EasyOCR 2.50, hybrid_pdf 2.27, Tesseract OCR 2.27, smart_pdf 2.27
Fails / empty	PyMuPDF 0.00, pdfminer 0.00, pdfplumber 0.00, MarkItDown 0.00, LangChain / LlamaIndex / Haystack loaders 0.00, OpenDataLoader 0.00

Tip: Docling is clearly ahead. hybrid_pdf and smart_pdf (Tesseract only) fall to about 2.3.

App pick: OK here. The app picks Docling for this type of file: Docling scores 3.89, within 0.1 of the best.

Tested on 5 documents: AI_agent_scanned_noisy, scanned_invoice_low_quality, scanned_invoice_skewed_noisy, scanned_financial_statement, scanned_purchase_order_phone.

Long scan (31 pages)

A report that is only photos of pages.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.58
Weak	PyMuPDF4LLM 3.67, Unstructured 3.02, hybrid_pdf 2.66, Tesseract OCR 2.66, smart_pdf 2.66
Fails / empty	pdfplumber 0.00, PyMuPDF 0.00, MarkItDown 0.00, OpenDataLoader 0.00

Tip: Only Docling scores well, but it takes about 6 minutes.

App pick: OK here. The app picks Docling for this type of file: Docling scores 4.58, within 0.1 of the best.

Tested on 1 document (one document only: a hint, not proof): scanned_report_30p.

Pages with text and pages that are scans

Some pages digital, some pictures.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.51, PyMuPDF4LLM 4.44
Also good	hybrid_pdf 3.95, smart_pdf 3.95
Weak	EasyOCR 3.56, Tesseract OCR 3.34, PyMuPDF 3.08, OpenDataLoader 3.06, pdfplumber 2.86, MarkItDown 2.86, Unstructured 2.80, LangChain / LlamaIndex / Haystack loaders 2.58, pdfminer 2.56

Tip: Docling or PyMuPDF4LLM. hybrid_pdf and smart_pdf lose the tables on the scanned pages.

App pick: OK here. The app picks PyMuPDF4LLM for this type of file: PyMuPDF4LLM scores 4.44, within 0.1 of the best.

Tested on 3 documents: AI_agent_mixed, mixed, mixed_digital_scanned.

Image files (png, jpg)

Receipt, ID card, screenshot, handwritten note.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.85
Also good	EasyOCR 4.67, Tesseract OCR 4.50, Unstructured 4.44

Tip: Docling is best; the OCR engines are close.

App pick: OK here. The app picks Docling for this type of file: Docling scores 4.85, within 0.1 of the best.

Tested on 4 documents: handwritten_note.png, id_card.png, receipt_photo.jpg, screenshot.png.

Office files

Word

Manuals, procedures, handbook (3 to 40 sections), lists and a merged table row.

Level	Parsers (average score, 0 to 5)
Best (equal)	Unstructured 4.68

Also good	Docling 4.49, native Word reader 4.40, MarkItDown 4.38

Tip: Unstructured and Docling are a little ahead of the native reader, which drops the page header and flattens list numbers.

App pick: Slightly below the best, on purpose. The app picks native Word reader for this type of file: native Word reader scores 4.40, 0.28 below the best (4.68). Kept first on purpose: another parser must win by more than 0.5 to replace the native reader.

Tested on 6 documents: procedure, report, report_with_figures, quality_manual, sop_goods_receipt, handbook_40p.

PowerPoint

Decks, with or without real chart objects.

Level	Parsers (average score, 0 to 5)
Best (equal)	Docling 4.62
Also good	MarkItDown 4.22, native PowerPoint reader 4.12
Weak	Unstructured 3.88

Tip: Native reader and Docling are both good. Docling reads real charts; the native reader cannot.

App pick: Slightly below the best, on purpose. The app picks native PowerPoint reader for this type of file: native PowerPoint reader scores 4.12, 0.50 below the best (4.62). Kept first on purpose: another parser must win by more than 0.5 to replace the native reader.

Tested on 4 documents: deck, slides, dashboard, quarterly_review.

Excel

Several sheets, formulas, merged title cells, a hidden column, charts, up to 1,500 rows.

Level	Parsers (average score, 0 to 5)
Best (equal)	native Excel reader 4.41, MarkItDown 4.37
Also good	Unstructured 4.27, Docling 3.91

Tip: The native reader is at or near the top. MarkItDown and Unstructured are about equal.

App pick: OK here. The app picks native Excel reader for this type of file: native Excel reader scores 4.41, within 0.1 of the best.

Tested on 6 documents: multi_sheet, sales, sales_with_chart, budget_model, hr_headcount_report, big_workbook.

CSV

Plain data files.

Level	Parsers (average score, 0 to 5)
Best (equal)	native Excel reader 5.00, MarkItDown 5.00, Unstructured 5.00

Tip: Every reader gets it right.

App pick: OK here. The app picks native Excel reader for this type of file: native Excel reader scores 5.00, within 0.1 of the best.

Tested on 2 documents: parts.csv, sales.csv.

Plain text and Markdown

Notes and runbooks.

Level	Parsers (average score, 0 to 5)
Best (equal)	plain text reader 5.00
Weak	Unstructured 2.50

Tip: Use the plain text reader only. Other parsers can fail on odd text files.

App pick: OK here. The app picks plain text reader for this type of file: plain text reader scores 5.00, within 0.1 of the best.

Tested on 2 documents: runbook.md, shift_notes.txt.

3. Good to know

- "App picks" is the first parser Smart Extract tries for that type of file. It chooses by file type, not by the situation inside the file, so it can miss (for example two tables side by side).
- hybrid_pdf = PyMuPDF + Tesseract OCR. smart_pdf = the same plus a vision model that reads charts and photos (max 8 per document, needs "Include paid parsers").
- adi_vision = Azure Document Intelligence reads text and tables, a vision model reads only charts and photos (max 8 per document; never a whole page). Paid, never picked automatically. It matches the best parsers on merged-header tables and scans, and is weaker on nested tables.
- Docling is often best but slow: minutes on documents of 30 to 100 pages.
- Some situations rest on one or two documents; the text says so. All documents were generated by us, not real company files. Differences under 0.1 are noise.
- Not tested: Azure Document Intelligence (paid). vlm (sends every page to the LLM) is switched off.